

Large-scale Distributed Systems


Lecture 2: Basic abstractions

Today

- Define **basic abstractions** that capture the fundamental characteristics of distributed systems.
 - We will later define more elaborate abstractions on top of those.
- Three main abstractions:
 - **Process** abstractions
 - **Link** abstractions
 - **Timing** abstractions
- A **distributed system model** = a combination of the three categories of abstractions.

Need for distributed abstractions

- Core of any distributed system is a **set of distributed algorithms**.
 - Implemented as a middleware between network (OS) and the application.
- **Reliable** applications need underlying services **stronger** than transport protocols (e.g., TCP or UDP).



Network protocols are not enough

- Communication
 - Reliability guarantees (e.g. with TCP) are only offered for **one-to-one** communication (client-server).
 - How to do **group communication**?
- High-level services
 - Sometimes one-to-many communication is not enough.
 - Need reliable **higher-level services**.
- Strategy: build complex distributed systems in a **bottom-up** fashion, from simpler ones.

High level services:


shared memory
consensus
atomic commit
group membership

Group communication:

reliable broadcast
causal order broadcast
total order broadcast
terminating reliable broadcast

Distributed computation

Distributed computation



- A **distributed algorithm** is a distributed collection $\Pi = \{p, q, r, \dots\}$ of N processes implemented by **identical** automata.
- The automaton at a process regulates the way the process executes its computation steps.
- Processes jointly implement the application.
 - Need for **coordination**.


Programming with events

- Every process consists of **modules** or **components**.
 - Modules may exist in multiple instances.
 - Every instance has a unique identifier and is characterized by a set of properties.
- Asynchronous **events** represent **communication** or **control flow** between components.
 - Each component is constructed as a state-machine whose transitions are triggered by the reception of events.
 - Events carry information (sender, message, etc)
- Reactive programming model:

```
    upon event  $\langle co_1, Event_1 \mid att_1^1, att_1^2, \dots \rangle$  do
        do something;
        trigger  $\langle co_2, Event_2 \mid att_2^1, att_2^2, \dots \rangle$ ;           // send some event

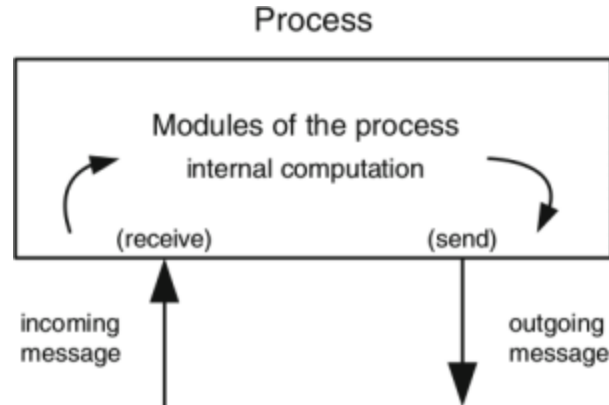
    upon event  $\langle co_1, Event_3 \mid att_3^1, att_3^2, \dots \rangle$  do
        do something else;
        trigger  $\langle co_2, Event_4 \mid att_4^1, att_4^2, \dots \rangle$ ;       // send some other event
```
- Effectively, a distributed algorithm is described by a set of event handlers.

Layered modular architecture



- Components can be composed locally to build software stacks.
 - The top of the stack is the [application layer](#).
 - The bottom of the stack the [transport](#) or [network](#) layer.
- Distributed programming abstraction layers are typically in the middle.
- We assume that every process executes the code triggered by events in a mutually exclusive way, without concurrently processing ≥ 2 events.

Execution



- The **execution** of a distributed algorithm is a **sequence of steps** executed by its processes.
- A **process step** consists in
 - **receiving** a message from another process,
 - **executing** a local computation,
 - **sending** a message to some process.
- Local messages between components are treated as local computation.
- We assume **deterministic** process steps (with respect to the message received and the local state prior to executing a step).

Liveness and safety

- Implementing a distributed programming abstraction requires satisfying its **correctness** in all possible executions of the algorithm.
 - i.e., in all possible interleaving of steps.
- Correctness of an abstraction is expressed in terms of **liveness** and **safety** properties.
 - **Safety**: properties that state that nothing bad ever happens.
 - A safety property is a property such that, whenever it is violated in some execution E of an algorithm, there is a prefix E' of E such that the property will be violated in any extension of E' .
 - **Liveness**: properties that state something good eventually happens.
 - A liveness property is a property such that for any prefix E' of E , there exists an extension of E' for which the property is satisfied.
- Any property can be expressed as the conjunction of safety property and a liveness property.

Correctness examples

Traffic lights at an intersection

- Safety: only one direction should have a green light.
- Liveness: every direction should eventually get a green light.



Correctness examples

TCP

- Safety: messages are not duplicated and received in the order they were sent.
- Liveness: messages are not lost.
 - i.e., messages are eventually delivered.

Assumptions

- In our abstraction of a distributed system, we need to specify the **assumptions** needed for the algorithm to be **correct**.
- A distributed system model includes assumptions on:
 - **failure** behavior of processes and channels
 - **timing** behavior of processes and channels

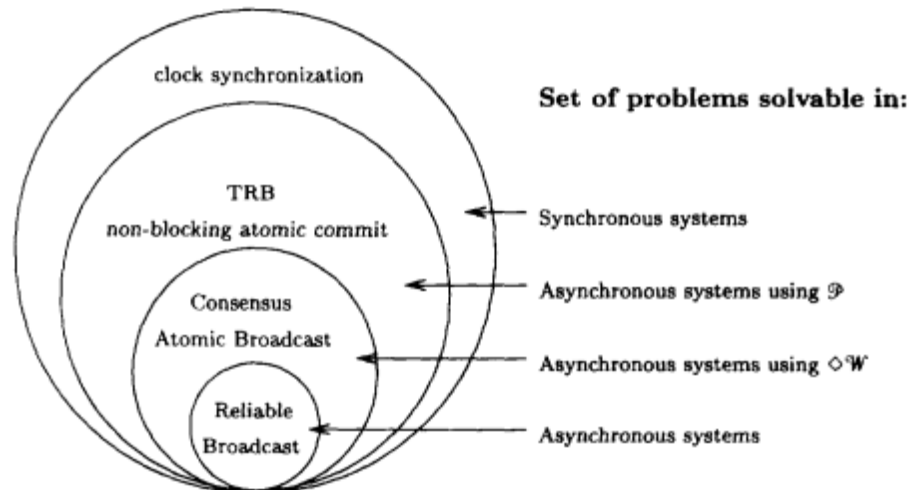


FIG. 9. Problem solvability in different distributed computing models.

Together, these assumptions define sets of solvable problems.

Process abstractions

Process failures

- Processes may **fail** in four different ways:
 - Crash-stop
 - Omissions
 - Crash-recovery
 - Byzantine / arbitrary
- Processes that do not fail in an execution are **correct**.

Crash-stop failures

- A process **stops taking steps**.
 - Not sending messages.
 - Not receiving messages.
- We assume the crash-stop process abstraction **by default**.
 - Hence, do not recover.
 - [Q] Does this mean that processes are not allowed to recover?

Omission failures

- Process **omits** sending or receiving messages.
 - **Send omission**: A process omits to send a message it has to send according to its algorithm.
 - **Receive omission**: A process fails to receive a message that was sent to it.
- Often, omission failures are due to **buffer overflows**.
- With omission failures, a process deviates from its algorithm by dropping messages that should have been exchanged with other processes.


Crash-recovery failures

- A process **might crash**.
 - It stops taking steps, not receiving and sending messages.
- It may **recover** after crashing.
 - The process emits a <Recovery> event upon recovery.
- Access to **stable storage**:
 - May read/write (**expensive**) to permanent storage device.
 - Storage survives crashes.
 - E.g., save state to storage, crash, recover, read saved state, ...
- A failure is different in the crash-recovery abstraction:
 - A process is **faulty** in an execution if
 - It crashes and never recovers, or
 - It crashes and recovers infinitely often.
 - Hence, a **correct** process may crash and recover.

Byzantine failures

- A process may **behave arbitrarily**.
 - Sending messages not specified by its algorithm.
 - Updating its state as not specified by its algorithm.
- Might behave **maliciously**, attacking the system.
 - Several malicious nodes might collude.

Fault-tolerance hierarchy




[Q] Explain how failure modes are special cases of one another.


Communication abstractions

Links


- Every process may **logically** communicate with every other process (a).
- The physical implementation may **differ** (b-d).




(a)



(b)



(c)



(d)

Link failures

- Fair-loss links
 - Channel delivers any message sent, with non-zero probability.
- Stubborn links
 - Channel delivers any message sent infinitely many times.
 - Can be implemented using fair-loss links.
- Perfect links (reliable)
 - Channel delivers any message sent exactly once.
 - Can be implemented using stubborn links.
 - **By default**, we assume the perfect links abstraction.
- [Q] What abstraction do UDP and TCP implement?

Stubborn links: interface

Module:

Name: StubbornPointToPointLinks, **instance** *sl*.

Events:

Request: $\langle sl, Send \mid q, m \rangle$: Requests to send message m to process q .

Indication: $\langle sl, Deliver \mid p, m \rangle$: Delivers message m sent by process p .

Properties:

SL1: *Stubborn delivery*: If a correct process p sends a message m once to a correct process q , then q delivers m an infinite number of times.

SL2: *No creation*: If some process q delivers a message m with sender p , then m was previously sent to q by process p .

[Q] Which property is safety/liveness/neither?

Perfect links: interface

Module:

Name: PerfectPointToPointLinks, **instance** *pl*.

Events:

Request: $\langle pl, Send \mid q, m \rangle$: Requests to send message *m* to process *q*.

Indication: $\langle pl, Deliver \mid p, m \rangle$: Delivers message *m* sent by process *p*.

Properties:

PL1: *Reliable delivery*: If a correct process *p* sends a message *m* to a correct process *q*, then *q* eventually delivers *m*.

PL2: *No duplication*: No message is delivered by a process more than once.

PL3: *No creation*: If some process *q* delivers a message *m* with sender *p*, then *m* was previously sent to *q* by process *p*.

[Q] Which property is safety/liveness/neither?

Perfect links: implementation

Implements:

PerfectPointToPointLinks, **instance** *pl*.

Uses:

StubbornPointToPointLinks, **instance** *sl*.

upon event $\langle pl, Init \rangle$ **do**

delivered := \emptyset ;

upon event $\langle pl, Send \mid q, m \rangle$ **do**

trigger $\langle sl, Send \mid q, m \rangle$;

upon event $\langle sl, Deliver \mid p, m \rangle$ **do**

if $m \notin delivered$ **then**

delivered := *delivered* $\cup \{m\}$;

trigger $\langle pl, Deliver \mid p, m \rangle$;

[Q] How does TCP efficiently maintain its delivered log?

Correctness of PL

- **PL1. Reliable delivery**
 - Guaranteed by the Stubborn link abstraction. (The Stubborn link will deliver the message an infinite number of times.)
- **PL2. No duplication**
 - Guaranteed by the log mechanism.
- **PL3. No creation**
 - Guaranteed by the Stubborn link abstraction.

Timing assumptions

Timing assumptions

- Timing assumptions correspond to the **behavior** of processes and links **with respect to the passage of time**. They relate to
 - different processing speeds of processes;
 - different speeds of messages (channels).
- Three basic types of system:
 - Asynchronous system
 - Synchronous system
 - Partially synchronous system

Asynchronous systems


- **No timing assumptions** on processes and links.
 - Processes do not have access to any sort of physical clock.
 - Processing time may vary arbitrarily.
 - No bound on transmission time.
- But **causality** between events can still be determined.
 - How?

Causal order

The **happened-before** relation $e_1 \rightarrow e_2$ denotes that e_1 may have caused e_2 . It is true in the following cases:

- **FIFO order**: e_1 and e_2 occurred at the same process p and e_1 occurred e_2 ;
- **Network order**: e_1 corresponds to the transmission of m at a process p and e_2 corresponds to its reception at a process q ;
- **Transitivity**: if $e_1 \rightarrow e'$ and $e' \rightarrow e_2$, then $e_1 \rightarrow e_2$.

Causal order



Similarity of executions

- The **view** of p in E , denoted $E|p$ is the subsequence of process steps in E restricted to those of p
- Two executions E and F are **similar w.r.t. to p** if $E|p = F|p$.
- Two executions E and F are **similar** if $E|p = F|p$ for all processes p .

Computation theorem


If two executions E and F have the same collection of events and their **causal order** is preserved, then E and F are similar executions.

Logical clocks

In an asynchronous distributed system, the passage of time can be measured with **logical clocks**:

- Each process has a local logical clock l_p , initially set a 0.
- Whenever an event occurs locally at p or when a process sends a message, p increments its logical clock.
 - $l_p := l_p + 1$
- When p sends a message event m , it timestamps the message with its current logical time, $t(m) := l_p$.
- When p receives a message event m with timestamp $t(m)$, p updates its logical clock.
 - $l_p := \max(l_p, t(m)) + 1$

Logical clocks



Clock consistency condition

Logical clocks capture **cause-effect relations**:

$$e_1 \rightarrow e_2 \Rightarrow t(e_1) < t(e_2)$$


- If e_1 is the cause of e_2 , then $t(e_1) < t(e_2)$.
 - Can you prove it?
- But not necessarily the opposite:
 - $t(e_1) < t(e_2)$ does not imply $e_1 \rightarrow e_2$.
 - e_1 and e_2 may be logically **concurrent**.

Vector clocks

Vector clocks fix this issue by making it possible to tell when two events cannot be causally related, i.e. when they are concurrent.

- Each process p maintains a vector V_p of N clocks, initially set at $V_p[i] = 0 \forall i$.
- Whenever an event occurs locally at p or when a process sends a message, p increments the p -th element of its vector clock.
 - $V_p[p] := V_p[p] + 1$
- When p sends a message event m , it piggybacks its vector clock as $V_m := V_p$.
- When p receives a message event m with the vector clock V_m , p updates its vector clock.
 - $V_p[p] := V_p[p] + 1$
 - $V_p[i] := \max(V_p[i], V_m[i]), \text{ for } i \neq p$.

Vector clocks



Comparing vector clocks

- $V_p = V_q$
 - iff $\forall i V_p[i] = V_q[i]$.
- $V_p \leq V_q$
 - iff $\forall i V_p[i] \leq V_q[i]$.
- $V_p < V_q$
 - iff $V_p \leq V_q$ AND $\exists j V_p[j] < V_q[j]$
- V_p and V_q are logically concurrent.
 - iff NOT $V_p \leq V_q$ AND NOT $V_q \leq V_p$

Synchronous systems

Assumption of three properties:

- Synchronous computation
 - Known upper bound on the process computation delay.
- Synchronous communication
 - Known upper bound on message transmission delay.
- Synchronous physical clocks
 - Processes have access to a local physical clock;
 - Known upper bound on clock drift and clock skew.

[Q] Why studying synchronous systems? What services can be provided?

Partially synchronous systems

A partially synchronous system is a system that is synchronous **most of the time**.

- There are periods where the timing assumptions of a synchronous system do not hold.
- But the distributed algorithm will have a long enough time window where everything behaves nicely, so that it can achieve its goal.

[Q] Are there such systems?

Timing abstractions

Failure detection

- It is **tedious** to model (partial) synchrony.
- Timing assumptions are mostly needed to detect failures.
 - Heartbeats, timeouts, etc.
- We define **failure detector** abstractions to **encapsulate timing assumptions**:
 - Black box giving suspicions regarding node failures;
 - Accuracy of suspicions depends on model strength.

Implementation of failure detectors

A typical implementation is the following:

- Periodically exchange **heartbeat** messages;
- **Timeout** based on **worst case** message round trip;
- If timeout, then **suspect** node;
- If reception of a message from a suspected node, **revise suspicion** and increase timeout.

Perfect detector: interface

Assuming a crash-stop process abstraction, the **perfect detector** encapsulates the timing assumptions of a **synchronous system**.

Module:

Name: PerfectFailureDetector, **instance** \mathcal{P} .

Events:

Indication: $\langle \mathcal{P}, \text{Crash} \mid p \rangle$: Detects that process p has crashed.

Properties:

PFD1: *Strong completeness:* Eventually, every process that crashes is permanently detected by every correct process.

PFD2: *Strong accuracy:* If a process p is detected by any process, then p has crashed.

[Q] Which property is safety/liveness/neither?

Perfect detector: implementation

Implements:

PerfectFailureDetector, **instance** \mathcal{P} .

Uses:

PerfectPointToPointLinks, **instance** pl .

upon event $\langle \mathcal{P}, \text{Init} \rangle$ **do**

$alive := \Pi$;

$detected := \emptyset$;

$starttimer(\Delta)$;

upon event $\langle \text{Timeout} \rangle$ **do**

forall $p \in \Pi$ **do**

if $(p \notin alive) \wedge (p \notin detected)$ **then**

$detected := detected \cup \{p\}$;

trigger $\langle \mathcal{P}, \text{Crash} \mid p \rangle$;

trigger $\langle pl, \text{Send} \mid p, [\text{HEARTBEATREQUEST}] \rangle$;

$alive := \emptyset$;

$starttimer(\Delta)$;

upon event $\langle pl, \text{Deliver} \mid q, [\text{HEARTBEATREQUEST}] \rangle$ **do**

trigger $\langle pl, \text{Send} \mid q, [\text{HEARTBEATREPLY}] \rangle$;

upon event $\langle pl, \text{Deliver} \mid p, [\text{HEARTBEATREPLY}] \rangle$ **do**

$alive := alive \cup \{p\}$;

Correctness

We assume a synchronous system:

- The transmission delay is bounded by some known constant.
- Local processing is negligible.
- The timeout delay Δ is chosen to be large enough such that
 - every process has enough time to send a heartbeat message to all,
 - every heartbeat message has enough time to be delivered,
 - the correct destination processes have enough time to process the heartbeat and to send a reply,
 - the replies have enough time to reach the original sender and to be processed.

Correctness:

- **PFD1. Strong completeness**
 - A crashed process p stops replying to heartbeat messages, and no process will deliver its messages. Every correct process will thus eventually detect the crash of p .
- **PFD2. Strong accuracy**
 - The crash of p is detected by some other process q only if q does not deliver a message from p before the timeout period.
 - This happens only if p has indeed crashed, because the algorithm makes sure p must have sent a message otherwise and the synchrony assumptions imply that the message should have been delivered before the timeout period.

Eventually perfect detector: interface

The **eventually perfect detector** encapsulates the timing assumptions of a **partially synchronous system**.

Module:

Name: EventuallyPerfectFailureDetector, **instance** $\Diamond\mathcal{P}$.

Events:

Indication: $\langle \Diamond\mathcal{P}, \text{Suspect} \mid p \rangle$: Notifies that process p is suspected to have crashed.

Indication: $\langle \Diamond\mathcal{P}, \text{Restore} \mid p \rangle$: Notifies that process p is not suspected anymore.

Properties:

EPFD1: *Strong completeness:* Eventually, every process that crashes is permanently suspected by every correct process.

EPFD2: *Eventual strong accuracy:* Eventually, no correct process is suspected by any correct process.

Eventually perfect detector: impl.

Implements:

EventuallyPerfectFailureDetector, instance $\diamond\mathcal{P}$.

Uses:

PerfectPointToPointLinks, instance pl .

upon event $\langle pl, Deliver \mid q, [HEARTBEATREQUEST] \rangle$ **do**
 trigger $\langle pl, Send \mid q, [HEARTBEATREPLY] \rangle$;

upon event $\langle pl, Deliver \mid p, [HEARTBEATREPLY] \rangle$ **do**
 $alive := alive \cup \{p\}$;

upon event $\langle \diamond\mathcal{P}, Init \rangle$ **do**

$alive := \Pi$;
 $suspected := \emptyset$;
 $delay := \Delta$;
 starttimer(delay);

upon event $\langle Timeout \rangle$ **do**

if $alive \cap suspected \neq \emptyset$ **then**

$delay := delay + \Delta$;

forall $p \in \Pi$ **do**

if $(p \notin alive) \wedge (p \notin suspected)$ **then**

$suspected := suspected \cup \{p\}$;

trigger $\langle \diamond\mathcal{P}, Suspect \mid p \rangle$;

else if $(p \in alive) \wedge (p \in suspected)$ **then**

$suspected := suspected \setminus \{p\}$;

trigger $\langle \diamond\mathcal{P}, Restore \mid p \rangle$;

trigger $\langle pl, Send \mid p, [HEARTBEATREQUEST] \rangle$;

$alive := \emptyset$;

 starttimer(delay);

[Q] Show that this implementation is correct.

Leader election

- Failure detection captures failure behavior.
 - Detects **failed** processes.
- **Leader election** is an abstraction that also captures failure behavior.
 - Detects **correct** nodes.
 - But a single and same for all, called the **leader**.
- If the current leader crashes, a new leader should be elected.

Leader election: interface

Module:

Name: LeaderElection, **instance** *le*.

Events:

Indication: $\langle le, Leader \mid p \rangle$: Indicates that process *p* is elected as leader.

Properties:

LE1: *Eventual detection*: Either there is no correct process, or some correct process is eventually elected as the leader.

LE2: *Accuracy*: If a process is leader, then all previously elected leaders have crashed.

Leader election: implementation

Implements:

LeaderElection, **instance** le .

Uses:

PerfectFailureDetector, **instance** \mathcal{P} .

upon event $\langle le, Init \rangle$ **do**

$suspected := \emptyset$;

$leader := \perp$;

upon event $\langle \mathcal{P}, Crash \mid p \rangle$ **do**

$suspected := suspected \cup \{p\}$;

upon $leader \neq \text{maxrank}(\Pi \setminus suspected)$ **do**

$leader := \text{maxrank}(\Pi \setminus suspected)$;

trigger $\langle le, Leader \mid leader \rangle$;

[Q] Show that this implementation is correct.

[Q] Is LE a failure detector?

Distributed system models

Distributed system models

We define a **distributed system model** as the combination of (i) a process abstraction, (ii) a link abstraction, and (iii) a failure detector abstraction.

- **Fail-stop** (synchronous)
 - Crash-stop process abstraction
 - Perfect links
 - Perfect failure detector
- **Fail-silent** (asynchronous)
 - Crash-stop process abstraction
 - Perfect links
- **Fail-noisy** (partially synchronous)
 - Crash-stop process abstraction
 - Perfect links
 - Eventually perfect failure detector
- **Fail-recovery**
 - Crash-stop process abstraction
 - Stubborn links

The fail-stop distributed system model substantially simplifies the design of distributed algorithms.

References

- Alpern, Bowen, and Fred B. Schneider. "Recognizing safety and liveness." Distributed computing 2.3 (1987): 117-126.
- Lamport, Leslie. "Time, clocks, and the ordering of events in a distributed system." Communications of the ACM 21.7 (1978): 558-565.
- Fidge, Colin J. "Timestamps in message-passing systems that preserve the partial ordering." (1987): 56-66.